



Making posters with Runpod Serverless

By Jessica Garson Beauchemin





What is Runpod?

Runpod Serverless

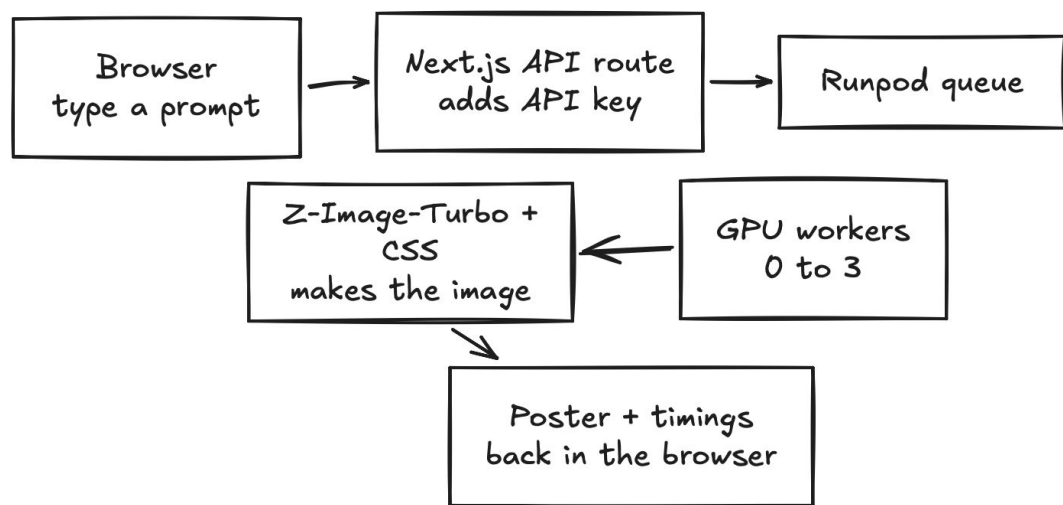
A cloud computing platform that lets you serve AI models for inference and run other compute-intensive workloads without managing servers. You only pay for the actual compute time you use, with no idle costs when your application isn't processing requests.



I built a live poster
generator that shows
how Serverless is
working in real-time

Architecture

At a high level your browser sends a prompt to the server, the server queues a job on Runpod, and the browser keeps asking "is it done yet?" until the image comes back and appears on a CSS-designed poster.



QUEUE CLEAR

Queue clear. Workers are winding down.

Workers stay up for the endpoint's idle timeout, then scale down to zero. That timeout is the dial between startup latency and idle cost.



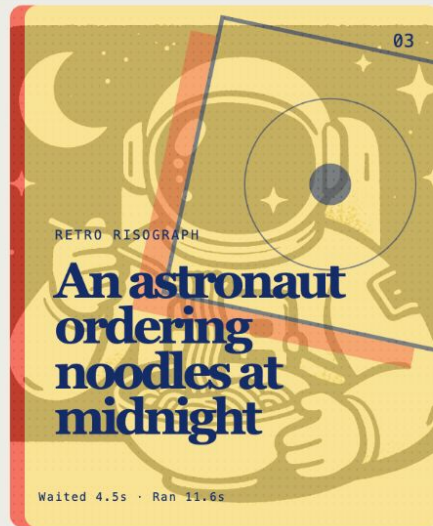
Where the request goes



0	1	28.3s
QUEUED	RUNNING	LAST JOB

● GPU WORKER 1	RUNNING
● GPU WORKER 2	IDLE
● GPU WORKER 3	IDLE

Live from the Runpod endpoint health API





If anything comes up along the way
or you want to connect with others
building with Runpod, join our
Discord

discord.com/invite/runpod

